

Operating System: Chap6 Process Synchronization

National Tsing Hua University
2021, Fall Semester



Overview

- Background
- Critical Section
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Thread Programming
- Atomic Transactions

Background

- **Concurrent** access to shared data may result in **data inconsistency**
- Maintaining data consistency requires mechanism to **ensure the orderly execution** of cooperating processes

Consumer & Producer Problem

■ Determine whether buffer is empty or full

- Previously: use *in*, *out* position
- Now: use *count* value

```
/*producer*/
```

```
while (1) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE) ;  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
/*consumer*/
```

```
while (1) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```

Concurrent Operations on *counter*

- The statement “*counter++*” may be implemented in machine language as:

move ax, counter

add ax, 1

move counter, ax

- The statement “*counter--*” may be implemented as:

move bx, counter

sub bx, 1

move counter, bx

Instruction Interleaving

- Assume counter is initially 5. One interleaving of statement is:

producer: move ax, counter → ax = 5

producer: add ax, 1 → ax = 6

context switch

consumer: move bx, counter → bx = 5

consumer: sub bx, 1 → bx = 4

context switch

producer: move counter, ax → counter = 6

context switch

consumer: move counter, bx → counter = 4

- The value of counter may be either 4, 5, or 6, where the correct result should be 5

Race Condition

- **Race condition**: the situation where several processes **access and manipulate shared data concurrently**. The final value of the shared data depends upon which process **finishes last**
- To prevent race condition, **concurrent** processes must be **synchronized**
 - On a single-processor machine, we could **disable interrupt** or use **non-preemptive CPU scheduling**
- Commonly described as **critical section problem**



Critical Section

The Critical-Section Problem

- *Purpose: a protocol for processes to cooperate*
- *Problem description:*
 - *N processes* are competing to use some *shared data*
 - Each process has a **code segment**, called *critical section*, in which the shared data is accessed
 - Ensure that when one process is executing in its critical section, *no other process is allowed to execute in its critical section* → *mutually exclusive*

The Critical-Section Problem

■ General code section structure

- Only one process can be in a critical section

do {

entry section

critical section

exit section

remainder section

} while (1);

Get entry permission

Modify shared data

Release entry permission

Critical Section Requirements

1. **Mutual Exclusion:** if process P is executing in its CS, no other processes can be executing in their CS
 2. **Progress:** if no process is executing in its CS and there exist some processes that wish to enter their CS, these processes cannot be postponed indefinitely
 3. **Bounded Waiting:** A bound must exist on the number of times that other processes are allowed to enter their CS after a process has made a request to enter its CS
- ➔ How to design entry and exist section to satisfy the above requirement?

Review Slides (1)

- Race condition?
- Critical-Section (CS) problem? 4 sections?
 - entry, CS, exit, remainder
- 3 requirements for solutions to CS problems?
 - mutual exclusion
 - progress
 - bounded waiting




Critical Section Solutions & Synchronization Tools

- Software Solution
- Synchronization Hardware
- Semaphore
- Monitor

Algorithm for Two Processes

- Only 2 processes, P_0 and P_1
- Shared variables
 - int **turn**; //initially turn = 0
 - $\text{turn} = i \Rightarrow P_i$ can enter its critical section



➔ Mutual exclusion? **Yes** Progress? **No**
Bounded-Wait? **Yes**

Peterson's Solution for Two Processes

■ Shared variables

- `int turn;` //initially `turn = 0`
- `turn = i` $\Rightarrow P_i$ can enter its critical section
- `boolean flag[2];` //initially `flag [0] = flag [1] = false`
- `flag [i] = true` $\Rightarrow P_i$ ready to enter its critical section

//Pi:

do {

```
flag[ i ] = TRUE;
turn = j ;
while (flag [ j ] &&
       turn == j ) ;
```

critical section

```
flag [ i ] = FALSE ;
```

remainder section

} while (1) ;

Enter CS when **either**:

1. a process gets its turn
2. the other process is not ready

Proof of Peterson's Solution

■ Mutual exclusion:

- If P_0 CS \rightarrow $\text{flag}[1] == \text{false} \mid \mid \text{turn} == 0$
- If P_1 CS \rightarrow $\text{flag}[0] == \text{false} \mid \mid \text{turn} == 1$

■ Assume both processes in CS \rightarrow $\text{flag}[0] == \text{flag}[1] == \text{true}$

\rightarrow $\text{turn} == 0$ for P_0 to enter, $\text{turn} == 1$ for P_1 to enter

- However, "turn" will be either 0 or 1 because its value will be set for both processes, but only one value will last
- **Therefore, P_0, P_1 can't in CS at the same time!**

```
/* process 0 */
do {
    flag[ 0 ] = TRUE;
    turn = 1 ;
    while (flag [ 1 ] && turn == 1 ) ;
    → critical section
    flag [ 0 ] = FALSE ;
    remainder section
} while (1) ;
```

```
/* process 1 */
do {
    flag[ 1 ] = TRUE;
    turn = 0 ;
    while (flag [ 0 ] && turn == 0 ) ;
    → critical section
    flag [ 1 ] = FALSE ;
    remainder section
} while (1) ;
```


Proof of Peterson's Solution

■ Progress (e.g., P_0 wishes to enter its CS):

(1) If P_1 is not ready \rightarrow flag[1] = false \rightarrow P_0 can enter

(2) If both are ready \rightarrow flag[0] == flag[1] == true

◆ If turn == 0 then P_0 enters, otherwise P_1 enters

➤ **Either cases, some waiting process can enter CS!**

```
/* process 0 */
do {
  flag[ 0 ] = TRUE;
  turn = 1 ;
  → while (flag [ 1 ] && turn == 1 ) ;
    critical section
  flag [ 0 ] = FALSE ;
    remainder section
} while (1) ;
```

```
/* process 1 */
do {
  flag[ 1 ] = TRUE;
  turn = 0 ;
  (2) → while (flag [ 0 ] && turn == 0 ) ;
    critical section
  (1) → flag [ 1 ] = FALSE ;
    remainder section
} while (1) ;
```

Proof of Peterson's Solution

- Bounded waiting (e.g., P_0 wishes to enter its CS):
 - (1) Once P_1 exits CS \rightarrow $\text{flag}[1] == \text{false} \rightarrow P_0$ can enter
 - (2) If P_1 exits CS && reset $\text{flag}[1] = \text{true}$
 $\rightarrow \text{turn} == 0$ (overwrite P_0 setting) $\rightarrow P_0$ can enter

➤ **P_0 won't wait indefinitely!**

```
/* process 0 */
do {
    flag[ 0 ] = TRUE;
    turn = 1 ;
    while (flag [ 1 ] && turn == 1 ) ;
    critical section
    flag [ 0 ] = FALSE ;
    remainder section
} while (1) ;
```

```
/* process 1 */
do {
    flag[ 1 ] = TRUE;
    (2) turn = 0 ;
    while (flag [ 0 ] && turn == 0 ) ;
    critical section
    (1) flag [ 1 ] = FALSE ;
    remainder section
} while (1) ;
```

Producer/Consumer Problem

■ Producer process

```
while (TRUE) {  
    entry-section( );  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
    computing();  
    exit-section( );  
}
```

■ Consumer process

```
while (TRUE) {  
    entry-section( );  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    computing();  
    exit-section( );  
}
```

➔ Incorrect: deadlock, if consumer enters the CS first.

Producer/Consumer Problem

■ Producer process

```
while (TRUE) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter++;  
    computing();  
    exit-section( );  
}
```

■ Consumer process

```
while (TRUE) {  
    while (counter == 0);  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter--;  
    computing();  
    exit-section( );  
}
```

➔ Correct but poor performance

Producer/Consumer Problem

■ Producer process

```
while (TRUE) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter++;  
    exit-section( );  
    computing();  
}
```

■ Consumer process

```
while (TRUE) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    entry-section( );  
    counter--;  
    exit-section( );  
    computing();  
}
```

➔ Correct & Maximize concurrent performance

Bakery Algorithm (n processes)

- Before enter its CS, each process receives a #
- Holder of the smallest # enters CS
- The numbering scheme always generates # in non-decreasing order; i.e., 1,2,3,3,4,5,5,5
- If processes P_i and P_j receive the same #, if $i < j$, then P_i is served first
- Notation:
 - $(a, b) < (c, d)$ if $a < c$ or if $a == c \ \&\& \ b < d$

Bakery Algorithm (n processes)

//Process i:

do {

 choosing [i] = TRUE ; ←

Get ticket

 num[i] = max(num[0], num[1], ..., num[n-1]) + 1;

 choosing [i] = FALSE ; ←

 for (j = 0; j < n; j++) {

 while (choosing [j]) ; ←

Cannot compare when
num is being modified

 while ((num[j] != 0) &&
 ((num[j], j) < (num[i], i))) ;

FCFS

 }

critical section

release
ticket

num[i] = 0 ;

remainder section

} while (1) ;

- Bounded-waiting because processes enter CS on a First-Come, First Served basis

Bakery Algorithm (n processes)

- Why cannot compare when num is being modified?
- Without locking...
 - 1. Let 5 be the current maximum number
 - 2. If P1 and P4 take number together, but P4 finishes before P1
 - ◆ **P1 = 0**; P4 = 6 → P4 will enter the CS
 - 3. After P1 takes the number
 - ◆ **P1 = P4 = 6** → **P1 will enter the CS as well!!!**
- With locking...
 - P4 will have to wait until P1 finish taking the number
 - **Both P1 & P4 will have the new number “6” before comparison**



Synchronization HW

Hardware Support

- The CS problem occurs because the modification of a shared variable may be **interrupted**
- If disable interrupts when in CS...
 - not feasible in multiprocessor machine
 - clock interrupts cannot fire in any machine
- HW support solution: **atomic instructions**
 - atomic: **as one uninterruptible unit**
 - examples: **TestAndSet(var)**, **Swap(a,b)**

Atomic TestAndSet()

```
boolean TestAndSet ( bool &lock) {  
    bool value = lock ;  
    lock = TRUE ;  
    return value ;  
}
```

execute atomically:
return the value of “lock”
and set “lock” to TRUE

Mutual exclusion? **Yes** Progress? **Yes** Bounded-Wait? **No!**

Shared data: boolean lock; //initially lock = FALSE;

do { // P0

while (TestAndSet (lock)) ;

critical section

lock = FALSE;

remainder section

} while (1) ;

do { // P1

while (TestAndSet (lock)) ;

critical section

lock = FALSE;

remainder section

} while (1) ;

obtain lock

release lock

Atomic Swap()

- Idea: enter CS if lock==false:

Shared data: boolean **lock**; //initially **lock = FALSE**;

do { // P0

```
key0 = TRUE;
while (key0 == TRUE)
    Swap (lock, key0) ;
```

critical section

```
lock = FALSE;
```

remainder section

} while (1) ;

do { // P1

```
key1 = TRUE;
while (key1 == TRUE)
    Swap (lock, key1) ;
```

critical section

```
lock = FALSE;
```

remainder section

} while (1) ;

Mutual exclusion? **Yes** Progress? **Yes** Bounded-Wait? **No!**

Review Slide (2)

- Use software solution to solve CS?
 - Peterson's and Bakery algorithms
- Use HW support to solve CS?
 - TestAndTest(), Swap()



Semaphores

Semaphore

- A tool to generalize the synchronization problem (easier to solve, but no guarantee for correctness)
- More specifically...
 - a record of how many units of a particular resource are available
 - ◆ If #record = 1 → binary semaphore, mutex lock
 - ◆ If #record > 1 → counting semaphore
 - accessed only through 2 *atomic* ops: **wait** & **signal**
- **Spinlock** implementation:
 - Semaphore is an integer variable

```
wait (S) {                               signal (S) {
    while (S <= 0) ;                      S++;
    S--;                                  }
}
```

busy waiting

POSIX Semaphore

- Semaphore is part of **POSIX** standard BUT it is **not** belonged to Pthread
 - It can be used with or **without** thread
- POSIX Semaphore routines:
 - **sem_init**(sem_t *sem, int pshared, unsigned int value)
 - **sem_wait**(sem_t *sem)
 - **sem_post**(sem_t *sem)
 - **sem_getvalue**(sem_t *sem, int *valp)
 - **sem_destory**(sem_t *sem)

Initial value of the semaphore

Current value of the semaphore

■ Example:

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem);
sem_wait(&sem);
    // critical section
sem_post(&sem);
sem_destory(&sem);
```


n -Process Critical Section Problem

- shared data:

semaphore mutex ; // initially mutex = 1

- Process P_i :

do {

wait (mutex) ; // pthread_mutex_lock(&mutex)

critical section

signal (mutex); // pthread_mutex_unlock(&mutex)

remainder section

} while (1) ;

Progress? Yes

Bounded waiting? Depends on the implementation of wait()

Non-busy waiting Implementation


■ Semaphore is data struct with a queue

- may use any queuing strategy (FIFO, FILO, etc)

```
typedef struct {  
    int value; // init to 0  
    struct process *L ;  
    // "PCB" queue  
} semaphore ;
```

E.g.,:

Value = -3



■ wait() and signal()

- use system calls: **block()** and **wakeup()**
- must be executed **atomically**

```
void wait (semaphore S) {  
    S.value--; // subtract first  
    if (S.value < 0) {  
        add this process to S.L ;  
        block( );  
    }  
}
```

```
void signal (semaphore S) {  
    S.value++;  
    if (S.value <= 0) {  
        remove a process P from S.L ;  
        wakeup(P);  
    }  
}
```

Atomic Operation

- How to ensure atomic wait & signal ops?
 - Single-processor: disable interrupts
 - Multi-processor:
 - ◆ HW support (e.g. Test-And-Set, Swap)
 - ◆ SW solution (Peterson's solution, Bakery algorithm)

Semaphore with Critical Section

```
void wait (semaphore S) {  
    entry-section( );  
    S.value--;  
    if (S.value < 0) {  
        add this process to S.L ;  
    }  
    exit-section( );  
    sleep( );  
}  
else {  
    exit-section( );  
}  
}
```

```
void signal (semaphore S) {  
    entry-section( );  
    S.value++;  
    if (S.value <= 0)  
        remove a process P from S.L;  
    exit-section( );  
    wakeup(P);  
}  
else {  
    exit-section( );  
}  
}
```

■ Busy waiting for entry-section()?

- limited to only the CS of wait & signal (~10 instructions)
 - ➔ very short period of time

Cooperation Synchronization

- P1 executes S1 ; P2 executes S2
 - S2 be executed only after S1 has completed
- Implementation:
 - shared var:

semaphore **sync** ; // initially sync = 0

P1:

S1 ;
signal (**sync**) ;

P2:

wait (**sync**) ;
S2 ;

A More Complicated Example

(Initially, all semaphores are 0)

begin

P₁: S₁; signal(**a**); signal(**b**);

P₂: wait(**a**); S₂; signal(**c**);

P₃: wait(**b**); S₃; signal(**d**);


P₄: wait(**c**); S₄; signal(**e**); signal(**f**);

P₅: wait(**e**); S₅; signal(**g**);

P₆: wait(**f**); wait(**d**); S₆; signal(**h**);


P₇: wait(**g**); wait(**h**); S₇;

end



Deadlocks & Starvation

- **Deadlocks:** 2 processes are waiting indefinitely for each other to release resources
- **Starvation:** example: LIFO queue in semaphore process queue



Review Slide (3)

- What's semaphore? 2 operations?
- What's busy-waiting (spinlock) semaphore?
- What's non-busy-waiting (non-spinlock) semaphore?
- How to ensure atomic wait & signal ops?
- Deadlock? starvation?



Classical Synchronization Problems

Listing & Purpose

- Purpose: used for testing newly proposed synchronization scheme
- Bounded-Buffer (Producer-Consumer) Problem
- Reader-Writers Problem
- Dining-Philosopher Problem

Bounded-Buffer Problem

- A pool of n buffers, each capable of holding one item
- Producer:
 - grab an empty buffer
 - place an item into the buffer
 - waits if no empty buffer is available
- Consumer:
 - grab a buffer and retracts the item
 - place the buffer back to the free pool
 - waits if all buffers are empty

Readers-Writers Problem

- A set of shared data objects
- A group of processes
 - reader processes (read shared objects)
 - writer processes (update shared objects)
 - **a writer process has exclusive access to a shared object**
- *Different variations involving priority*
 - *first RW problem*: no reader will be kept waiting unless a writer is updating a shared object
 - *second RW problem*: once a writer is ready, it performs the updates as soon as the shared object is released
 - ➔ writer has higher priority than reader
 - ➔ once a writer is ready, no new reader may start reading

First Reader-Writer Algorithm

```
// mutual exclusion for write
semaphore wrt=1
// mutual exclusion for readcount
semaphore mutex=1
int readcount=0;
```

```
Writer(){
    while(TRUE){
        wait(wrt);
        // Writer Code

        signal(wrt);
    }
}
```

Acquire write lock
if **reads** haven't

```
Reader(){
    while(TRUE){
        wait(mutex);
        readcount++;
        if(readcount==1)
            wait(wrt);
        signal(mutex);
        // Reader Code


        wait(mutex);
        readcount--;
        if(readcount==0)
            signal(wrt);
        signal(mutex);
    }
}
```

release write lock if
no more reads

- Readers share a single wrt lock
- **Writer may have starvation problem**

Dining-Philosophers Problem

- 5 persons sitting on 5 chairs with 5 chopsticks
- A person is either thinking or eating
 - thinking: no interaction with the rest 4 persons
 - eating: need 2 chopsticks at hand
 - a person picks up 1 chopstick at a time
 - done eating: put down both chopsticks
- **deadlock problem**
 - one chopstick as one semaphore
- **starvation problem**





Monitors

Motivation

- Although semaphores provide a convenient and effective synchronization mechanism, its correctness is depending on the programmer
 - All processes access a shared data object must execute `wait()` and `signal()` in the right order and right place
 - This may not be true because honest programming error or uncooperative programmer

Monitor --- A high-level language construct

- The representation of a **monitor type** consists of
 - declarations of **variables** whose values define the state of an instance of the type
 - **Procedures/functions** that implement operations on the type
- The monitor type is similar to a **class in O.O. language**
 - A procedure within a monitor can access only **local variables** and the formal **parameters**
 - The local variables of a monitor can be used only by the local procedures
- But, the monitor ensures that **only one process at a time can be active** within the monitor
- Similar idea is incorporated to many prog. language:
 - concurrent pascal, C# and Java


Monitor

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes

Syntax

```
monitor monitor-name {  
    // shared variable declarations  
    procedure body P1 (...) {  
        ...  
    }  
    procedure body P2 (...) {  
        ...  
    }  
    procedure body Pn (...) {  
        ...  
    }  
    initialization code {  
    }  
}
```

Schematic View



Monitor Condition Variables

- To allow a process to **wait within** the monitor, a condition variable must be declared, as

condition x, y;

- Condition variable can only be used with the operations **wait()** and **signal()**

- **x.wait();**


- means that the process invoking this operation is suspended until another process invokes

- **x.signal();**

- resumes exactly one suspended process. If no process is suspended, then the signal operation **has no effect**

- (In contrast, signal always change the state of a semaphore)

Monitor With Condition Variables



Dining Philosophers Example

```
monitor dp {  
    enum {thinking, hungry, eating} state[5]; //current state  
    condition self[5]; //delay eating if can't obtain chopsticks  
    void pickup(int i)        // pickup chopsticks  
    void putdown(int i)      // putdown chopsticks  
    void test(int i)         // try to eat  
    void init() {  
        for (int i = 0; i < 5; i++)  
            state[i] = thinking;  
    }  
}
```

```

void pickup(int i) {
    state[i] = hungry;
    test(i); //try to eat
    if (state[i] != eating)
        self[i].wait(); //wait to eat
}

```

```

void putdown(int i) {
    state[i] = thinking;
    // check if neighbors
    // are waiting to eat
    test((i+4) % 5);
    test((i+1) % 5);
}

```

//try to let P_i eat (if it is hungry)


```

void test(int i) {
    if ( (state[(i + 4) % 5] != eating) &&(state[(i + 1) % 5] != eating)
        && (state[i] == hungry) ) {
        //No neighbors are eating and  $P_i$  is hungry
        state[i] = eating;
        self[i].signal(); ←
    }
}

```

If P_i is suspended, resume it
 If P_i is not suspended, **no effect**


An illustration



P1:
DiningPhilosophers.pickup(1)
eat
DiningPhilosophers.putdown(1)

P2:
DiningPhilosophers.pickup(2)
eat
DiningPhilosophers.putdown(2)

An illustration



```
void pickup(int i) {  
    state[i] = hungry;  
    test(i); //try to eat  
    if (state[i] != eating)  
        self[i].wait(); //wait to eat  
}
```


P1:

```
→ DiningPhilosophers.pickup(1)  
    eat  
DiningPhilosophers.putdown(1)
```

P2:

```
DiningPhilosophers.pickup(2)  
    eat  
DiningPhilosophers.putdown(2)
```


An illustration



```
void pickup(int i) {  
    state[i] = hungry;  
    test(i); //try to eat  
    if (state[i] != eating)  
        self[i].wait(); //wait to eat  
}
```


P1:

```
→ DiningPhilosophers.pickup(1)  
    eat  
DiningPhilosophers.putdown(1)
```


P2:

```
DiningPhilosophers.pickup(2)  
    eat  
DiningPhilosophers.putdown(2)
```

An illustration



An illustration




```
void putdown(int i) {  
    state[i] = thinking;  
    // check if neighbors  
    // are waiting to eat  
    test((i+4) % 5); test((i+1) % 5);  
}
```

P1:
DiningPhilosophers.pickup(1)
eat
→ DiningPhilosophers.putdown(1)

P2:
→ DiningPhilosophers.pickup(2)
eat
DiningPhilosophers.putdown(2)

An illustration




```
void putdown(int i) {
    state[i] = thinking;
    // check if neighbors
    // are waiting to eat
    test((i+4) % 5); test((i+1) % 5);
}
```

P1:
 DiningPhilosophers.pickup(1)
 eat
 → DiningPhilosophers.putdown(1)

P2:
 → DiningPhilosophers.pickup(2)
 eat
 DiningPhilosophers.putdown(2)

An illustration



P1:
DiningPhilosophers.pickup(1)
eat
→ DiningPhilosophers.putdown(1)

P2:
DiningPhilosophers.pickup(2)
eat
→ DiningPhilosophers.putdown(2)



Thread Programming

Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of type `pthread_mutex_t` and initialized with `pthread_mutex_init()`
- A mutex is destroyed with `pthread_mutex_destory()`
- A critical section can then be protected using `pthread_mutex_lock()` and `pthread_mutex_unlock()`
- Example:

```
#include "pthread.h"
pthread_mutex_t mutex;
pthread_mutex_init (&mutex, NULL);
pthread_mutex_lock(&mutex);
// enter critical section

Critical Section

pthread_mutex_unlock(&mutex);
pthread_mutex_destory(&mutex);
// leave critical section
```

specify default attribute for the mutex

Condition Variables (CV)

- CV represent some **condition** that a thread can:
 - Wait on, until the condition occurs; or
 - Notify other waiting threads that the condition has occurred
- Three operations on condition variables:
 - **wait()** --- **Block** until another thread calls **signal()** or **broadcast()** on the CV
 - **signal()** --- Wake up **one thread** waiting on the CV
 - **broadcast()** --- Wake up **all threads** waiting on the CV
- In Pthread, CV **type** is a **pthread_cond_t**
 - Use **pthread_cond_init()** to initialize
 - **pthread_cond_wait (&theCV, &somelock)**
 - **pthread_cond_signal (&theCV)**
 - **pthread_cond_broadcast (&theCV)**

Using Condition Variable

■ Example:

- A threads is designed to **take action when x=0**
- Another thread is responsible for decrementing the counter

```
pthread_cond_t  cond;  
pthread_cond_init (cond, NULL);
```

```
pthread_mutex_t  mutex;  
pthread_mutex_init (mutex, NULL);
```

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- All condition variable operation **MUST** be performed while a **mutex is locked!!!**

Using Condition Variable

```
action() {  
→ pthread_mutex_lock (&mutex)  
  while (x != 0)  
    pthread_cond_wait (cond, mutex);  
  pthread_mutex_unlock (&mutex);  
  take_action();  
}
```

```
→ counter() {  
  pthread_mutex_lock (&mutex)  
  x--;  
  if (x==0)  
    pthread_cond_signal (cond);  
  pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
2. Wait()
 1. Put the thread into sleep & releases the lock

1. Lock mutex

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

2. Wait()

1. Put the thread into **sleep & releases the lock**

1. **Waked up**, but the **thread is locked**

1. Lock mutex

2. Signal()

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

2. Wait()

1. Put the thread into **sleep & releases the lock**

1. **Waked up**, but the **thread is locked**

2. **Re-acquire lock** and **resume execution**

1. Lock mutex

2. Signal()

3. Releases the lock

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

2. Wait()

1. Put the thread into **sleep & releases the lock**

1. **Waked up**, but the **thread is locked**

2. **Re-acquire lock** and **resume execution**

3. Release the lock

1. Lock mutex

2. Signal()

3. Releases the lock

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    while (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

2. Wait()

1. Put the thread into sleep &

releases the lock

1. Waked up, but the thread is locked

2. **Re-acquire lock** and resume execution

3. Release the lock

1. Lock mutex

2. Signal()

3. Releases the lock

Another reason why
condition variable op.
MUST within mutex lock

ThreadPool Implementation

Task structure

```
typedef struct {  
    void (*function)(void *);  
    void *argument;  
} threadpool_task_t;
```

Threadpool structure

```
struct threadpool_t {  
    pthread_mutex_t lock;  
    pthread_cond_t notify;  
    pthread_t *threads;  
    threadpool_task_t *queue;  
    int thread_count;  
    int queue_size;  
    int head;  
    int tail;  
    int count;  
    int shutdown;  
    int started;  
};
```

Allocate thread and task queue

```
/* Allocate thread and task queue */  
pool->threads = (pthread_t *) malloc(sizeof(pthread_t) * thread_count);  
pool->queue = (threadpool_task_t *) malloc(sizeof(threadpool_task_t) * queue_size);
```


ThreadPool Implementation

```
static void *threadpool_thread(void *threadpool) ← thread handler  
function  
{  
    threadpool_t *pool = (threadpool_t *)threadpool;  
    threadpool_task_t task;  
  
    for(;;) {  
        /* Lock must be taken to wait on conditional variable */  
        pthread_mutex_lock(&(pool->lock));  
  
        /* Wait on condition variable, check for spurious wakeups.  
         * When returning from pthread_cond_wait(), we own the lock. */  
        while((pool->count == 0) && (!pool->shutdown)) {  
            pthread_cond_wait(&(pool->notify), &(pool->lock));  
        }  
    }  
}
```

ThreadPool Implementation

```
/* Grab our task */
task.function = pool->queue[pool->head].function;
task.argument = pool->queue[pool->head].argument;
pool->head += 1;
pool->head = (pool->head == pool->queue_size) ? 0 : pool->head;
pool->count -= 1;

/* Unlock */
pthread_mutex_unlock(&(pool->lock));

/* Get to work */
(*(task.function))(task.argument);
}
```

Synchronized Tools in JAVA

■ Synchronized Methods (Monitor)

- Synchronized method uses the method receiver as a lock
- Two invocations of synchronized **methods cannot interleave on the same object**
- When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object block until the first thread exits the object

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() { c++; }  
    public synchronized void decrement() { c--; }  
    public synchronized int value() { return c; }  
}
```

Synchronized Tools in JAVA

■ Synchronized Statement (Mutex Lock)

- Synchronized blocks uses the **expression** as a lock
- A synchronized Statement can only be executed once the thread has obtained a **lock for the object or the class that has been referred to in the statement**
- useful for improving concurrency **with fine-grained**

```
public void run()
{
    synchronized(p1)
    {
        int i = 10; // statement without locking requirement
        p1.display(s1);
    }
}
```

Review Slides (4)

- Bounded-buffer problem?
- Reader-Writer problem?
- Dining Philosopher problem?
- What is monitor and why need monitor?



Atomic Transactions

System Model

- **Transaction**: a collection of instructions (or instructions) that performs a single logic function
- **Atomic Transaction**: operations happen as a single logical unit of work, in its entirety, or not at all
- Atomic transaction is particular a concern for database system
 - Strong interest to use DB techniques in OS

File I/O Example

- Transaction is a series of **read** and **write** operations
- Terminated by **commit** (transaction successful) or **abort** (transaction failed) operation
- Aborted transaction must be **rolled back** to undo any changes it performed
 - It is part of the responsibility of the system to ensure this property

Log-Based Recovery

- **Record** to stable storage information about all **modifications by a transaction**
 - **Stable storage**: never lost its stored data
- **Write-ahead logging**: Each log record describes single **transaction write operation**
 - Transaction name
 - Data item name
 - Old & new values
 - Special events: $\langle T_i \text{ starts} \rangle$, $\langle T_i \text{ commits} \rangle$
- Log is used to **reconstruct the state of the data items** modified by the transactions
 - Use **undo** (T_i), **redo**(T_i) to recover data

Checkpoints

- When failure occurs, must consult the log to **determine which transactions must be re-done**
 - Searching process is time consuming
 - Redone may not be necessary for all transactions
- Use **checkpoints** to reduce the above overhead:
 - Output all **log records** to stable storage
 - Output all **modified data** to stable storage
 - Output a log record **<checkpoint>** to stable storage

Review Slides (5)

- What is atomic transaction?
- Purpose of commit, abort, rolled-back?
- How to use log and checkpoints?

Textbook Problem Set

- 6.1: The first known correct software solution to the critical-section problem for two processes was developed by Dekker. The two processes, P0 and P1, share the following variables:

```
boolean flag[2];/*initiallyfalse*/  
int turn;
```

The structure of process P_i ($i == 0$ or 1) is shown in the figure below; the other process is P_j ($j == 1$ or 0). Prove that the algorithm satisfies all three requirements for the critical-section problem.

```
do {  
    flag[i] = true;  
  
    while (flag[j]) {  
        if (turn == j) {  
            flag[i] = false;  
            while (turn == j)  
                ; // do nothing  
            flag[i] = true;  
        }  
    }  
  
    // critical section  
  
    turn = j;  
    flag[i] = false;  
  
    // remainder section  
} while (true);
```


Figure 6.42 The structure of process P_i in Dekker's algorithm.

Textbook Problem Set

- 6.2: Explain why disable interrupts are not appropriate for implementing synchronization primitives in multiprocessor systems.
- 6.5: A file is to be shared among different processes, each of which has a unique number. The file can be accessed simultaneously by several processes, subject to the following constraint: the sum of all unique numbers associated with all the processes currently accessing the file must be less than n . Write a monitor to coordinate access to the file.
- 6.9: Servers can be designed to limit the number of open connections. For example, a server may wish to have only N socket connections at any point in time. As soon as N connections are made, the server will not accept another incoming connection until an existing connection is released. Explain how semaphores can be used by a server to limit the number of concurrent connections.
- 6.12: Show how to implement the `wait()` and `signal()` semaphore operations in multiprocessor environments using the `TestAndSet()` instruction. The solution should exhibit minimal busy waiting.



Backup



Case Study: Solaris 2 Windows XP

Solaris 2 Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses **adaptive mutexes** for efficiency when protecting data from **short code** segments.
 - Mutex and semaphore always **serialize data accesses**
- Uses **condition variables** and **readers-writers** locks when **longer sections** of code need access to data.
 - Efficient for data that is **accessed frequently, but in a read-only manner**

Solaris 2 Adaptive Mutex

■ Multiprocessor system

➤ Data locked (i.e. in use)

- ◆ Locking thread is running → requesting thread spins on the mutex (**spinlock**)
- ◆ Locking thread is not in run state → requesting thread blocks on the mutex (**waiting lock**)

■ Uniprocessor system

➤ Requesting thread always blocks

Solaris 2 Turnstile

- Uses **turnstiles** to order the list of threads waiting to acquire either an **adaptive mutex** or **reader-writer lock**
 - A turnstile is a queue structure containing threads blocked on a lock
- To prevent a **priority inversion**, turnstiles are organized according to a ***priority-inheritance protocol***
 - Temporarily inherit the priority of the high-priority thread (blocked on this lock)



XP Synchronization

- Use **interrupt masks** to protect access to global resources on **uniprocessor systems** (disable interrupt)
- Uses **spinlocks** on **multiprocessor system**
- Dispatcher objects: either in signaled or nonsignaled state
 - Signaled: object is available immediately
 - Nonsignaled: object is not available
 - Thread queue associated with each object
 - *WaitForSingleObject* or *WaitForMultipleObjects*